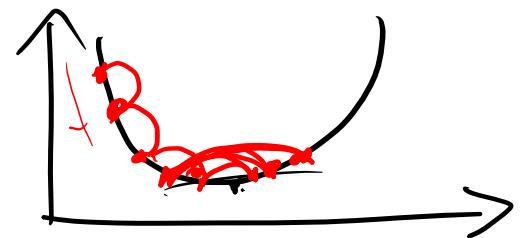
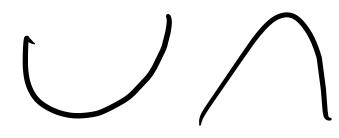


Gradient Descent

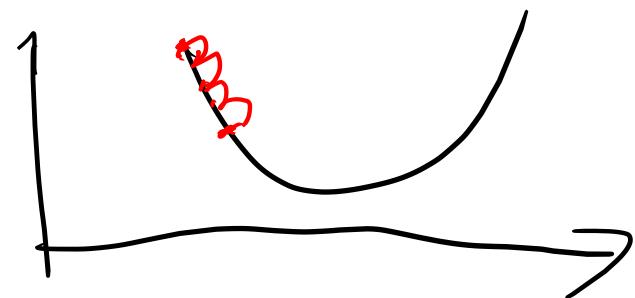


Normally: find derivative, set to zero,
analyse inflection point

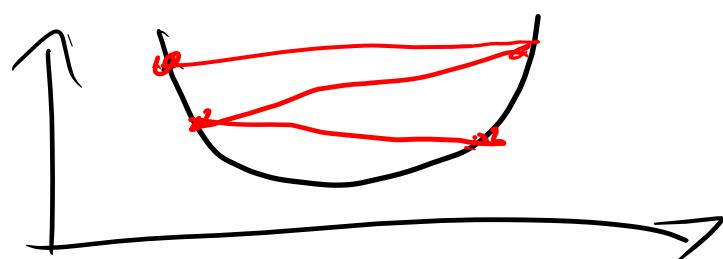


GD: iteratively descend in the direction
where change is greatest.

GD doesn't terminate! We need a stop-condition.
Step-length is called "learning rate".



If we use fixed number epochs
a too small lr will miss the optimum.

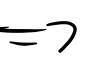


A too large lr will "skip over" the minimum.

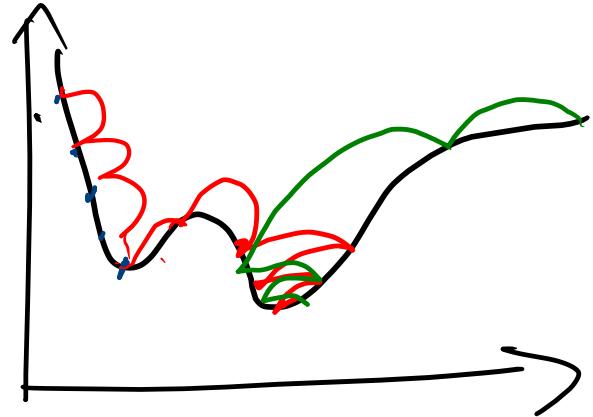
So . . . we want dynamic number of epochs
and adjust lr as we go.



Early Stopping



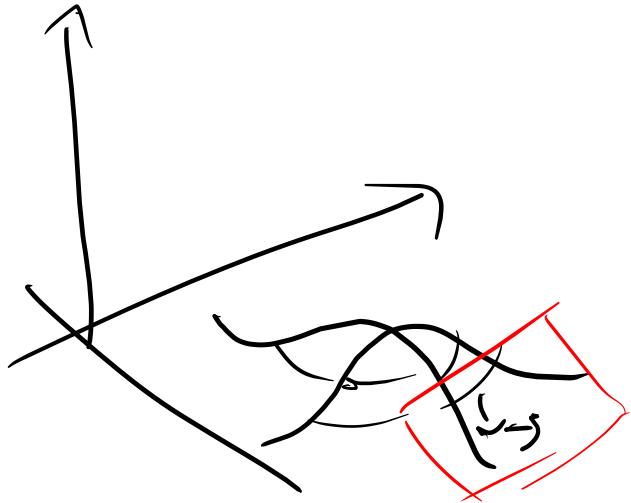
keep track of loss metric
when below threshold: stop



Adaptive Descent with Momentum (ADAM)

Adaptive learning rate : decreases with epoch

Momentum : we add a physics-like condition
to the cost function, that
really acts like momentum (friction)



∇f - gradient always points in the direction of largest change

- If we evaluate all points in the data:

Batch Gradient Descent

- expensive
- + good results

One step:

$$\theta_{n+1} = \theta_n - \eta \nabla_{\theta} \text{MSE}(\theta)$$

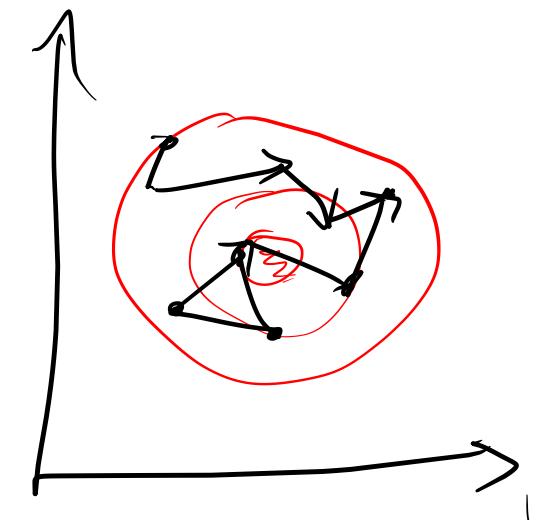
evaluated at each point for each parameter

Too expensive!

Stochastic Gradient Descent:

May or may not work: it's random!

- Choose a random point, derive at that point, take a small step.



Mini-batch Gradient Descent

Choose a random sample of points.

- evaluate, descend

By far the most common!